

Bhavya

Generative AI Engineer

varmabhavya38@gmail.com | +1 9197471470 | [LinkedIn](#)

PROFESSIONAL SUMMARY

AI Engineer with 3 years of experience building scalable Generative AI solutions using Large Language Models (LLMs), Retrieval-Augmented Generation (RAG), and Agentic AI architectures. Experienced in designing multi-agent systems with planner, executor, and verifier roles. Proficient in building end-to-end AI applications using LangGraph, LangChain, Hugging Face, FastAPI, and Streamlit, with a focus on performance, efficiency, and cost optimization through fine-tuning and model optimization.

TECHNICAL SKILLS

Programming Languages: Python

Generative AI & ML: LLMs (GPT-4, Claude, Gemini), Agentic AI, RAG, Prompt Engineering, Embeddings, Semantic Search, LLM Evaluation, Model Fine Tuning, NLP

Frameworks & Libraries: LangGraph, LangChain, LangSmith, Scikit-Learn, NLTK, Hugging Face Transformers

Vector Databases: FAISS, Pinecone, ChromaDB

Tools & Platforms: FastAPI, Streamlit, Git, Visual Studio Code, Excel VBA

EXPERIENCE

Cognizant Technology Solutions

Sep 2022 – Apr 2025

AI Engineer - Bengaluru, Karnataka

- Built a natural language SQL query assistant enabling non-technical users to query databases conversationally, export results to Excel, and auto-generate reports via email reducing manual reporting effort by ~60% and saving an estimated 8+ hours/week per team.
- Engineered a context-aware conversational AI system with persistent memory and vector-based retrieval using LangChain and Pinecone, improving grounded response accuracy and supporting context aware multi-turn interactions across 500+ daily queries.
- Built an Agentic AI orchestration framework using LangGraph where multiple autonomous agents (planner, executor, verifier) collaboratively executed multi step tasks reducing task failure rate compared to single agent baselines.
- Implemented an Agentic RAG system spanning multiple vector databases (Pinecone, FAISS, ChromaDB), improving retrieval precision and end to end answer accuracy on domain specific QA benchmarks.
- Designed LLM evaluation and monitoring pipelines using LangSmith and custom retrieval relevance metrics, reducing hallucination rate and enabling structured regression testing across model upgrades.
- Implemented end-to-end AI workflows combining LangChain, Hugging Face, FastAPI and Streamlit delivering interactive AI tools used by 3+ internal teams with real-time data processing.

Cognizant Technology Solutions

Mar 2022 – Sep 2022

AI Intern

- Developed Generative AI applications using Python, LLM APIs, and backend frameworks to create intelligent automation solutions.
- Built Retrieval-Augmented Generation (RAG) pipelines by implementing document ingestion, text chunking, embeddings, and context-based response generation.
- Integrated vector databases and similarity search techniques to improve information retrieval and enhance LLM-generated responses.
- Designed and tested prompt templates, evaluated model outputs, and optimized AI workflows for accuracy and relevance.
- Developed API-based AI services and collaborated with engineering teams using Git, debugging, and Agile development practices.

EDUCATION

Master of Science - Computer Science and Engineering

University of North Texas, Denton, TX

Bachelor of Technology - Computer Science and Engineering

Sri Venkateswara College of Engineering, Tirupati, Andhra Pradesh, India